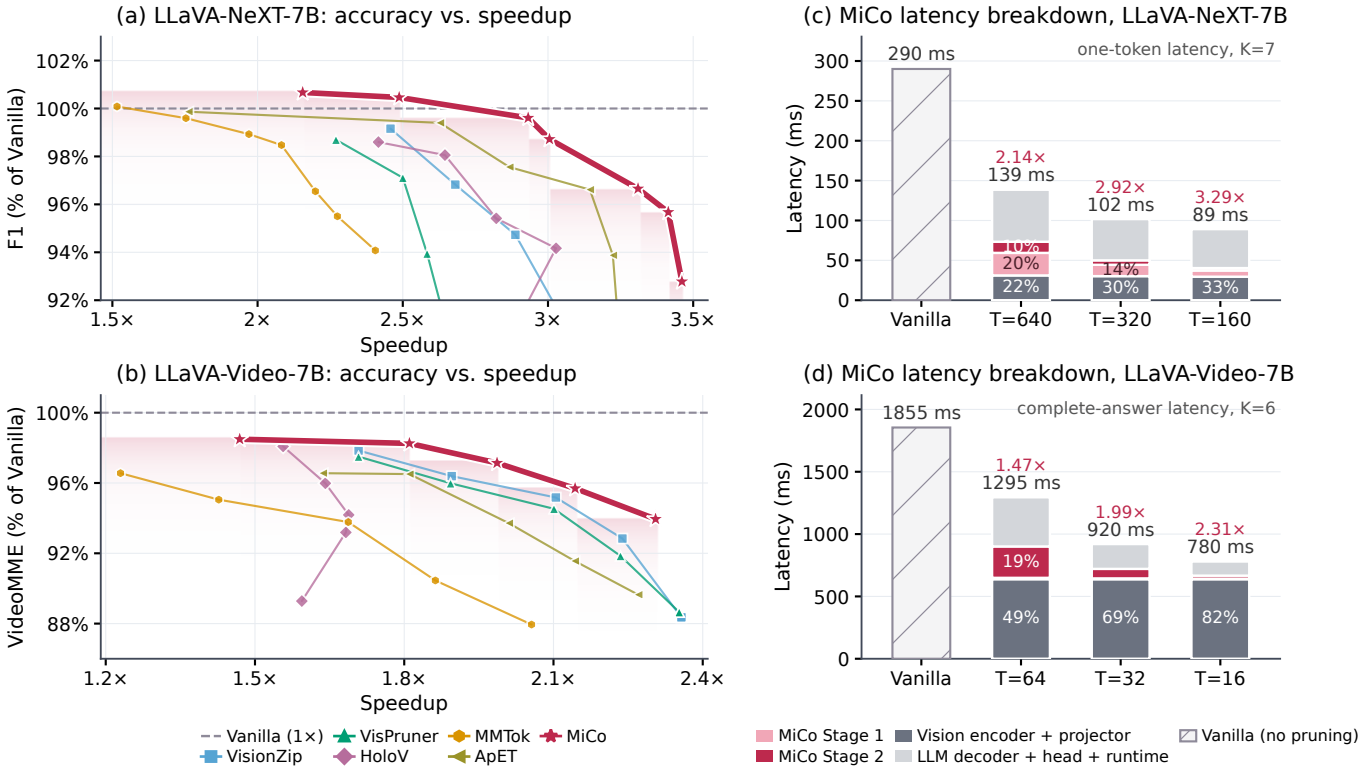
(a) LLaVA-NeXT-7B: accuracy vs. speedup
F1 (% of Vanilla)
102%
100%
98%
96%
94%
92%
1.5×
2×
2.5×
3×
3.5×
Speedup
(b) LLaVA-Video-7B: accuracy vs. speedup
VideoMME (% of Vanilla)
100%
96%
92%
88%
1.2×
1.5×
1.8×
2.1×
2.4×
Speedup
Vanilla (1×)
VisPruner
MMTok
MiCo
VisionZip
HoloV
ApET
(c) MiCo latency breakdown, LLaVA-NeXT-7B
one-token latency, K=7
Latency (ms)
300
250
200
150
100
50
0
290 ms
2.14×
139 ms
2.92×
102 ms
3.29×
89 ms
10%
20%
22%
14%
30%
33%
Vanilla
T=640
T=320
T=160
(d) MiCo latency breakdown, LLaVA-Video-7B
complete-answer latency, K=6
Latency (ms)
2000
1500
1000
500
0
1855 ms
1.47×
1295 ms
1.99×
920 ms
2.31×
780 ms
19%
49%
69%
82%
Vanilla
T=64
T=32
T=16
MiCo Stage 1
MiCo Stage 2
Vision encoder + projector
LLM decoder + head + runtime
Vanilla (no pruning)